

ClauseCheck: Predatory Debt Auditor & Regulatory Intent Classifier

Voice-Enabled Compliance, Speech Recognition & Deep Learning Pipeline for Consumer Credit Agreements

Done by: S. K. Thyakeshwar | Register No: 23BAI0194

Vellore Institute of Technology (VIT) • Email: thyakeshwar.sk2023@vitstudent.ac.in

Assessment: Lab Assessment • Total Marks: 20 • Domain: Fintech Law & Conversational AI • Repository: clausecheck/ • Status: Fully Operational & Validated

1. Executive Summary & Problem Formulation

In recent years, the rapid proliferation of digital lending applications and unregulated lending service providers (LSPs) across India has led to severe consumer exploitation. Predatory practices commonly include **undisclosed upfront deductions**, **non-transparent Annual Percentage Rates (APR)** disguised behind flat daily interest metrics, **unlawful scraping of mobile device contacts and private photo galleries**, **aggressive debt recovery harassment outside statutory legal hours**, **denial of mandated cooling-off exit windows**, and **unsolicited credit card issuance with compounded penalty charges**.


While the **Reserve Bank of India (RBI)** and the **Digital Personal Data Protection (DPDP) Act** have established comprehensive protective frameworks—most notably the *RBI Digital Lending Directions (2022/2024)*, the *Key Fact Statement (KFS) Mandate*, the *Fair Practices Code*, and the *Master Direction on Credit Cards*—everyday retail borrowers frequently lack the specialized legal literacy needed to audit dense loan agreements. Furthermore, vulnerable borrowers often express grievances through conversational voice queries containing regional slang, colloquial terms, or fragmented phrases.

ClauseCheck bridges this critical gap by delivering a voice-enabled, full-duplex regulatory compliance assistant. Borrowers can simply speak their contract clause or loan dispute into a microphone. The system captures the acoustic stream, transcribes the speech to text, computes dense semantic embeddings via a Transformer encoder, classifies the regulatory intent using a regularized deep PyTorch Multi-Layer Perceptron (MLP), verifies the primary and secondary risk vectors, renders an actionable statutory audit card with remediation rights, and synthesizes an audible voice response using Text-to-Speech (TTS).

Full-Duplex Voice Interaction Captures microphone audio via <code>st.audio_input</code> with robust PCM 16-bit transcoding via <code>soundfile</code> , transcribes speech via Google STT, and returns an audible synthesized voice verdict via <code>gTTS</code> .	Transformer Semantic Vectors Transforms raw user text into 384-dimensional dense semantic representations using <code>all-MiniLM-L6-v2</code> , capturing contextual meaning across formal terms and colloquialisms.
Deep PyTorch MLP Classifier A 3-layer feedforward network with Batch Normalization (<code>BatchNorm1d</code>) and Dropout ($p=0.3$) achieving 97.30% validation accuracy and 0.9736 Macro F1 across 6 regulatory intents.	Strict Hallucination Prevention Enforces dual-threshold out-of-domain abstention (<0.40 confidence or <code>out_of_scope</code> tag). Safely declines non-credit queries to prevent false legal advice.

2. System Architecture & Methodology

The ClauseCheck system operates as an integrated five-stage end-to-end processing pipeline, moving seamlessly from acoustic voice capture to statutory remedy synthesis:



Full Duplex Voice Pipeline: Acoustic Ingestion → SBERT Semantic Vectorization → Regularized PyTorch Classifier → Statutory RBI Card & Voice Feedback

Regulatory Anchors: RBI Digital Lending Directions (2022/2024), Fair Practices Code, Master Directions on Credit Cards, and DPDP Act Data Minimization

2.1 Audio Ingestion & Acoustic Preprocessing

User speech is captured directly in the browser using Streamlit's native `st.audio_input`. In modern browsers, audio buffers can arrive in diverse container formats (such as WebM or raw WAV) that can cause decoding failures in standard SpeechRecognition libraries. ClauseCheck resolves this through a robust fallback transcoding pathway using `soundfile`: the ingested byte stream is decoded into a NumPy floating-point array and transcoded to a standardized 16-bit linear PCM WAV buffer before being ingested by `speech_recognition.Recognizer`. A direct text input fallback is provided to support accessibility and environments where microphone permissions are restricted.

2.2 Semantic Vectorization & Neural Network Topology

Text queries are transformed into fixed-length 384-dimensional dense vectors using the pre-trained Sentence-Transformers `all-MiniLM-L6-v2` model. These vectors are subsequently classified by a custom PyTorch deep neural network structured as follows:


Layer	Layer Type	Dimensions / Parameters	Activation & Regularization
Input	Dense Feature Vector	384 features (Sentence-BERT)	L2 Normalized
Hidden 1	Linear + BatchNorm1d	Linear(384 → 64), 24,576 weights	ReLU, Dropout (p = 0.3)
Hidden 2	Linear	Linear(64 → 32), 2,048 weights	ReLU, Dropout (p = 0.3)
Output	Linear + Softmax	Linear(32 → 6), 192 weights	Softmax Probability Vector (6 classes)

Training Hyperparameters: The model is trained using the AdamW optimizer with a learning rate of 0.005 and weight decay of 0.01. A CosineAnnealingLR scheduler (`T_max=100`, `eta_min=1e-4`) smoothly decays the learning rate across 100 epochs. Model selection is governed by a strict checkpointing criterion that tracks held-out validation Macro F1 score, saving the optimal state dictionary to `models/intent_model.pth`.

3. Dataset Engineering & Regulatory Anchors

The training dataset (data/intents.json) consists of **370 curated and augmented samples** distributed across 6 distinct intent categories. To ensure exceptional resilience against automatic speech recognition (ASR) phonetic noise, colloquial Indian English, and fragmented syntax, the baseline 50 samples per class were expanded through domain-specific syntactic perturbations and speech transcription error models.

Dataset Composition Across Intent Classes (Total: 370 Samples)



Intent Tag	Count	Statutory / Regulatory Anchor	Sample Utterance & Domain Coverage
apr_hidden_fees	65	RBI Digital Lending Guidelines (2022/2024) & Key Fact Statement (KFS) Mandate	'Why is the lending app deducting a 1500 rupee processing fee upfront before disbursal?'
coercive_device_permissions	63	RBI Restriction on Mobile Device Data Access & DPDP Act Data Minimization	'The loan app refuses to approve my application unless I grant access to my contacts and photo gallery.'
recovery_agent_harassment	62	RBI Fair Practices Code & Recovery Norms (Calling hours 8 AM - 7 PM, anti-intimidation)	'A recovery agent is calling my relatives and threatening police arrest for a delayed payment.'
cooling_off_cancellation	62	Statutory Look-up / Cooling-Off Period Norms (1 to 3 days exit without foreclosure fine)	'I accepted this digital loan by mistake two days ago, can I return the money without penalty?'
credit_card_unilateral_terms	61	RBI Master Direction on Credit & Debit Cards (7-day closure, unsolicited limit protection)	'The bank increased my credit card limit without my consent and delayed closing my card.'
out_of_scope	57	Negative Baseline / Out-of-Domain Control Vector	'What is the capital city of Australia and how do I bake sourdough bread from scratch?'

4. Empirical Results & Quantitative Evaluation

Evaluation was conducted on a held-out test partition of **74 samples (20% stratified split)** unseen during gradient optimization. The model achieved state-of-the-art classification performance with zero training divergence.

Validation Accuracy 97.30% (72 / 74 test samples)	Macro-Averaged F1 0.9736 Balanced across all 6 classes	Weighted-Averaged F1 0.9734 Normalized sample weight	Average Inference Latency 15.1 ms Real-time CPU execution
--	---	---	--





Figure: (Left) Loss trajectory and held-out validation accuracy across 100 epochs; (Right) 6x6 multiclass evaluation confusion matrix demonstrating high diagonal concentration.

5. Multi-Scenario Stress Evaluation (28 Test Vectors)

To prove true generalization beyond the training distribution and ensure real-world readiness for diverse consumer inputs, the pipeline was subjected to a rigorous **28-case stress benchmark** spanning 7 distinct challenge categories.



Category	Cases	Pass Rate	Mean Latency	Behavioral Insight & Stress Condition
Standard In-Domain	5	100%	40.1 ms	Perfect classification across canonical phrasing and regulatory references.
Colloquial & Regional Slang	5	100%	22.5 ms	Robust to Indian colloquialisms like 'hafta vasooli', 'cut money', and 'yaar'.
ASR Phonetic Noise & Typos	5	100%	24.3 ms	Invariant to severe transcription errors: 'prosessing fee before desbursement', 'culling of peryod'.
Formal Contract Excerpts	3	100%	28.2 ms	Accurately parses dense legal language: 'irrevocably grants unconditional access to contact book'.
Multi-Violation Hybrid	1	100%	22.4 ms	Surfaces primary vector (apr_hidden_fees) and flags secondary correlated breach.
Minimal Keywords	5	100%	15.1 ms	Instant categorization of short 2-3 word queries (e.g., 'recovery agent abuse').
Out-of-Scope Negative Controls	4	100%	39.8 ms	Zero hallucination: safely abstains on cooking recipes, coding tasks, and random noise.
Overall Stress Benchmark	28 / 28	100.0%	27.5 ms	Flawless Generalization & Production Reliability

6. Regulatory Compliance Audit Engine & Statutory Remediation

Unlike generic chatbots that return conversational opinions, ClauseCheck maps every detected intent directly to enforceable statutory provisions under Indian banking law. When a violation is identified, the system generates a structured regulatory audit card comprising the exact statutory anchor, violation severity, predatory mechanism analysis, and legal remediation pathways:

1. Hidden Fees & APR Non-Transparency (apr_hidden_fees)

- **Statutory Anchor:** RBI Digital Lending Guidelines (2022/2024) & KFS Mandate.
- **Verdict:** HIGH RISK • Violation of Standardized KFS Requirements.
- **Legal Mechanism:** Lenders must disclose an all-inclusive Annual Percentage Rate (APR) before disbursal. Deducting upfront processing charges or third-party fees not itemized in the KFS is strictly prohibited.
- **Remedy:** File a formal complaint with the lender's Principal Nodal Officer and escalate to the RBI Complaint Management System (cms.rbi.org.in).

2. Unlawful Access to Phone Contacts & Storage (coercive_device_permissions)

- **Statutory Anchor:** RBI Digital Lending Norms & DPDP Act Data Minimization Principles.
- **Verdict:** CRITICAL VIOLATION • Prohibited Data Harvesting Vector.
- **Legal Mechanism:** Digital Lending Apps (DLAs) are explicitly barred from accessing mobile contact books, photo galleries, call records, and local file storage. Permissions are restricted to one-time KYC camera/mic use.
- **Remedy:** Demand immediate deletion of harvested telemetry under the DPDP Act and report the app to the state Cyber Crime Cell.

3. Coercive Recovery Tactics & Intimidation (recovery_agent_harassment)

- **Statutory Anchor:** RBI Guidelines on Fair Practices Code & Recovery Agent Conduct.
- **Verdict:** CRITICAL VIOLATION • Criminal Intimidation & Harassment Breach.
- **Legal Mechanism:** Agents cannot call before 8:00 AM or after 7:00 PM, contact family/friends, visit homes unannounced, or threaten arrest. Regulated entities remain strictly liable for agent misconduct.
- **Remedy:** Lodge a criminal intimidation FIR at the local police station and file an actionable grievance with the RBI Ombudsman.

4. Denial of Statutory Loan Exit Rights (cooling_off_cancellation)

- **Statutory Anchor:** RBI Mandated Look-up / Cooling-Off Period for Digital Loans.
- **Verdict:** REGULATORY NON-COMPLIANCE • Cooling-Off Provision Breach.
- **Legal Mechanism:** Borrowers possess an explicit statutory right to exit a digital loan without foreclosure penalties by repaying the principal plus proportionate APR within the look-up window (min. 3 days for ≥ 7 -day tenures).
- **Remedy:** Submit written notice invoking statutory look-up rights under RBI Digital Lending Guidelines with proof of tender.

5. Credit Card Unilateral Limit & Closure Delays (credit_card_unilateral_terms)

- **Statutory Anchor:** RBI Master Direction – Credit Card and Debit Card Issuance and Conduct Directions.
- **Verdict:** HIGH RISK • Violation of Credit Card Governance Norms.
- **Legal Mechanism:** Card issuers cannot unilaterally enhance limits or dispatch unsolicited cards. Card closure must be completed within 7 working days, failing which statutory compensation is triggered.
- **Remedy:** Demand statutory delay compensation of ₹500 per day directly from the issuing bank for each day past 7 days.

7. Voice Synthesis & User Interface Runtime

The front-end user experience is implemented in app.py via Streamlit, offering an intuitive, high-speed interface suitable for desktop and mobile web environments:

- **Two-Way Full-Duplex Voice:** Ingests spoken queries via st.audio_input, transcribes via Google STT, and plays back an automated voice verdict generated using gTTS synthesized with an Indian English acoustic profile (tld='co.in').
- **Dual-Vector Risk Surfacing:** Identifies primary intent and flags correlated secondary violations whenever secondary confidence exceeds 15% (e.g., hybrid predatory agreements combining hidden fees with contact scraping).
- **Probability Distribution Transparency:** An interactive expander provides a complete breakdown of multi-class softmax probabilities with animated progress meters.
- **Production Resource Caching:** PyTorch weights and the SBERT transformer are cached in memory using @st.cache_resource, keeping query processing latency under 25 milliseconds per invocation.
- **Zero Hallucination Abstinence:** When confidence is below 0.40 or classified as out_of_scope, the UI explicitly abstains and instructs the user rather than fabricating statutory conclusions.

8. Conclusion & Future Roadmap

ClauseCheck demonstrates that combining pre-trained dense Transformer embeddings with a lightweight, regularized deep PyTorch classification network yields an exceptionally fast, highly accurate (**97.30% accuracy, 100% stress benchmark pass rate**), and robust compliance auditor for predatory debt agreements.

Future Engineering Directions:

1. **On-Device Whisper ASR:** Transitioning from cloud-based Google STT to local quantised Whisper (e.g., whisper-tiny.en) to enable private, offline execution without cloud dependency.
2. **Indic Multilingual Support:** Extending semantic representations using paraphrase-multilingual-MiniLM-L12-v2 to support native Hindi, Tamil, Telugu, and Marathi voice complaints.
3. **Automated Legal Notice Generation:** Integrating dynamic PDF complaint generation that prepopulates formal RBI Ombudsman grievance letters with timestamps and violation clauses.

9. Regulatory & Technical References

1. Reserve Bank of India. *Guidelines on Digital Lending*. RBI/2022-23/111, September 2022.
2. Reserve Bank of India. *Master Direction – Reserve Bank of India (Credit Card and Debit Card – Issuance and Conduct) Directions, 2022*. RBI/2022-23/92, April 2022.
3. Ministry of Law and Justice, Government of India. *The Digital Personal Data Protection Act, 2023*. Act No. 22 of 2023.
4. Reimers, N., & Gurevych, I. (2019). *Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks*. In Proceedings of EMNLP 2019.
5. Paszke, A., et al. (2019). *PyTorch: An Imperative Style, High-Performance Deep Learning Library*. Advances in Neural Information Processing Systems (NeurIPS 32).